



ElectroCom response to the LORiMDT Clinical Governance Platform RFI

Submitted to: Casi M. Vician, Founder & CEO, Lucas Health Tech (casi@lucashealthtech.com)

Submitted by: Arslan Noor, CTO, ElectroCom (ElectroCom Innovations LLC),
arslan@electrocomit.com

Date: July 31, 2026

Confidential. Prepared for Lucas Health Tech under the terms of the LORiMDT RFI.

Cover note

Casi,

Thank you for including ElectroCom, and for the FAQ. It made the intent behind each requirement clear, and we wrote this response to it. We follow your numbering (questions 7 to 41), stay inside the 20-page limit, and recommend a tier with our reasoning rather than bidding to a number.

Two things up front, since you said an honest read is worth more than a polished bid. First, we recommend Tier 1 (Bootstrap MVP) with a pre-agreed phased upgrade to Tier 2 on your August funding close. We think that is the tier your requirements and funding reality call for, and we explain why in Section 2.8. Second, as I mentioned over email, healthcare is where my team works day to day: for the past two years we have modernized a legacy clinical platform into a live, cloud-based patient-management system now running across hospitals in Australia and the UK. We keep the client unnamed in this written response and can walk you through the engagement in full, including the client, at the demo stage.

Where LORiMDT asks for something we would be building for the first time, we say so plainly. And because we would rather you test us than take our word, we have proposed a short paid discovery sprint (Section 2.7, Q36) so you can watch us work before the full build commits.

Arslan

2.1 Company background and healthcare experience

7. Company overview.

ElectroCom trades as ElectroCom (legal entity ElectroCom Innovations LLC), a US-registered company founded in 1998 with 28 years in technology, from telecom infrastructure into full-stack software, data, and AI/ML delivery. Headquarters: 2750 S Preston Rd, Ste 1167, Celina, TX 75009. Team: 11+ across engineering, data science, and DevOps, of whom 6 have worked on our current healthcare SaaS engagement. We are privately held and services-profitable; we are not disclosing revenue figures at the RFI stage. Our PHI controls and data-residency approach are in Section 2.4.

8. Healthcare client base.

Our healthcare work is focused rather than broad, and we would rather show depth on one comparable build than a long logo wall. Our primary engagement is a multi-year modernization of a legacy clinical platform into a live, cloud-based patient-management SaaS, now serving hospitals across Australia and the UK. We keep the client unnamed in this written response and will walk you through the engagement in full, including the client, at the demo stage. We have also built a patient-management system for a healthcare-sector client in Pakistan; see references in Section 2.9.

9. Oncology experience.

Our directly comparable experience is the clinical-platform engagement above. It began as a modernization of a legacy desktop clinical application that had run hospital lab workflows for over a decade, and we rebuilt it module by module into a production React cloud application while the platform stayed live for its hospital users. It now runs across hospitals in Australia and the UK and covers patient records, clinical quality control, structured clinical reporting and diagnostic-report analysis, appointment scheduling, and staff rostering, on real clinical data. Our data team has hands-on experience turning clinical reports into structured, analyzable data, which maps directly to LORiMDT's registry-quality capture. The program ran about 20 months and never took the clinical service offline, which is the same constraint LORiMDT carries: you cannot pause a live clinical service to modernize it. See Appendix A for a fuller write-up.

We are being straight about the gap. That platform is not oncology-specific, and we have not built a tumor-board product before. What transfers directly is the hard part of LORiMDT's Tier 1: multi-hospital operation on real clinical data, longitudinal structured capture generated by the workflow itself, integration against legacy and standards-based systems, multi-tenant isolation, and an audit-grade governance layer. The HCC-specific clinical content (the BCLC-CUSE mapping, Child-Pugh and MELD-Na gating, LI-RADS capture, and the guideline-anchored thresholds) we would build to your validated field dictionary and under the clinical direction of Dr. Talia Baker and

your advisory group, the model your FAQ Q25 and Q51 describe. We are not proposing to author clinical logic; we are proposing to build the platform that runs yours.

10. Startup and early-stage experience.

This is our core client type. We work with seed to Series B companies that have a real product to ship and a team too small to execute the roadmap in-house, and we structure engagements around milestone-based, funding-aware terms. For LORiMDT that means we can start the Tier 1 build against a committed minimum for the August to October sprints (your FAQ Q19) and expand as your capital close lands, rather than requiring the full budget committed on day one. We are comfortable with the shared-risk, phased structure you signaled.

2.2 Technical capabilities

11. Platform architecture.

Cloud-native and container-based, deployed to your own US AWS environment as a managed, HIPAA-eligible footprint, which is the "on-premises to the customer's cloud" model your FAQ Q46 describes. Multi-region (East and West) for high availability and disaster recovery, with a day-one availability target around 98% maturing toward the five-nines objective by general availability (your FAQ Q46). We commit to an RPO of 15 minutes or less and an RTO of 4 hours or less at pilot, tightening as the platform matures, and we will align to the uniform RPO/RTO you publish for all vendors. We architect for secure multi-tenancy from day one, per your FAQ Q39: a single logical deployment across the 2 to 3 pilot sites, each an isolated tenant with strict logical data separation, site-scoped access control, and per-site configuration, rather than a separate stack per site. Proposed stack: containerized services on ECS or EKS, PostgreSQL for the longitudinal patient-journey data model, a FHIR integration layer, and infrastructure-as-code for reproducible per-site connection configuration and tenant isolation.

12. FHIR R4 support.

We build native HL7 FHIR R4 support. For LORiMDT specifically:

- Bidirectional at Tier 1 on the Epic connection: broad reads (Patient, Condition, MedicationRequest, Observation, DiagnosticReport, Procedure, and notes) and a narrow, governed write for the MDT note (see Section 2.3 and your FAQ Q28 and Q36).
- SMART on FHIR with EHR launch as the target Epic access pattern, distributed through App Orchard, so LORiMDT opens in context with patient and encounter established. A standalone backend and API mode is architected as the portable fallback and for non-Epic sites (your FAQ Q27).

- US Core profiles as the interoperability baseline, and mCODE as the target oncology profile for the cancer data model, extended with LORiMDT's governance elements (the clinical handshake, BCLC-CUSE mapping, liver-function gating, and the re-entry loop) where mCODE does not reach (your FAQ Q24 and Q38).
- A three-tier capture model as the design premise, because not all 850+ elements are FHIR-populable (your FAQ Q34): Auto-EHR from FHIR where the resource exists, AI-assisted drafting for clinician confirmation, and clinician-attested entry where no discrete source exists, with an HL7 v2 fallback at non-Epic sites.

13. EHR integration.

Epic and Oracle Health / Cerner via FHIR R4, with an HL7 v2 fallback where a resource is not exposed (your FAQ Q32). We design for separate, independent Epic installs across sites, which is the harder and more general case, with per-site connection configuration and tenant isolation, so anything that works across independent installs also works if two sites share a Community Connect instance (your FAQ Q29). Integration methods: FHIR REST APIs as the primary path, SMART on FHIR launch, HL7 v2 as the fallback, and CDA where needed.

On App Orchard specifically: we are not currently a listed Epic App Orchard developer. Our existing healthcare platform integrates with legacy and standards-based systems rather than Epic. We would complete Epic App Orchard registration and the certification required for the encounter-scoped MDT-note write as part of the Tier 1 build, and we budget for it. Your FAQ Q16 already frames the pilot as first-time integration bring-up at the sites, so this is planned work, not a surprise.

14. Workflow configuration.

We deliver a configurable clinical workflow engine, delivered configured to LORiMDT's defined pilot workflows, with your clinical-governance roles able to adjust configuration afterward. Not hardcoded screens, and not a bare toolkit you configure alone (your FAQ Q44). The core pilot workflow is the case lifecycle: initiation and data population, asynchronous specialist pre-review, synchronous MDT deliberation, recommendation and treatment-plan capture, the clinical handshake (ownership assignment plus acknowledgment), note generation and write-back, and the iterative reassessment and re-entry loop. It is conditional and includes stop points: required-field enforcement gates submission, per-cancer-type conditional fields drive visibility, and the Level 1 soft-stop halts progression until the accountable clinician acknowledges. Access is role- and site-scoped (submitting clinician, MDT chair, specialist reviewer, coordinator).

15. Rules engine and trigger logic.

We build a configurable, versioned rules engine, not hardcoded logic. It supports triggers on lab values, imaging and LI-RADS results, staging criteria, and time-based escalations. Rule changes are governed, versioned, and audit-logged (who changed what, when, and why), managed by

your authorized clinical-governance roles (your FAQ Q26). Cases in flight are pinned to the rule version under which they were being evaluated, so a mid-stream change does not retroactively rewrite a decision the board already acted on. Where a change is clinically material, the platform flags affected in-flight cases for re-review rather than mutating them. On ownership (your FAQ Q25): we build the engine, Lucas Health Tech owns the rule definitions and governance framework, and the clinical thresholds are guideline-anchored (AASLD, NCCN, ESMO, BCLC) and specialist-validated, not vendor-authored.

16. AI/ML capabilities.

This is our deepest strength. Our CTO, Arslan, spent the last year working with OpenAI on training their video-generation models, so the AI side is real rather than a wrapper, and our team ships production LLM and RAG systems. For LORiMDT, AI is applied narrowly, with a clinician always in control:

- Automated note generation at Tier 1. The platform assembles a structured MDT note from the case's own structured data plus the board's decision and drafts it into a formatted note. This is AI-assisted drafting, not autonomous authorship. The accountable clinician reviews it, edits the narrative and the discrete recommendation and plan elements, and only then commits it through the handshake (your FAQ Q33 and Q35).
- AI-assisted capture at Tier 1: drafting field values from available data for clinician confirmation, inside the three-tier model.
- Later tiers: NLP extraction from narrative radiology reports (Tier 3, not a Tier 1 promise, per your FAQ Q30), treatment-sequencing suggestions, trial-eligibility screening, and guideline-concordance checking with deviation flags.

We do not put AI in the decision path. It reduces transcription and drafting burden. The clinician owns every committed record.

17. API and extensibility.

REST APIs across platform services, webhook and event support, and an internal orchestration engine that can optionally consume FHIR Subscriptions where a site exposes them (for example, to refresh data when new results post) but does not require them. Tier 1 orchestration is user-action-driven through the platform's own state machine (your FAQ Q37). Third-party integrations such as Open Evidence at Tier 3 are designed as pluggable capabilities.

2.3 Clinical workflow and governance

18. MDT / tumor board support.

Yes. The case is modeled as a persistent, versioned, collaboratively edited object with explicit state, not a document that exists only during a live meeting (your FAQ Q22). We support the hybrid board: asynchronous preparation (case assembly, data population, specialist pre-review over days) followed by a synchronous sign-off event where the recommendation is deliberated and the clinical handshake (ownership assignment plus acknowledgment) is executed. Fully asynchronous virtual boards are supported and confirmed per site during discovery. The full audit trail of decisions, recommendations, and acknowledgments is captured (see Q21).

19. Role-based views.

Yes. Role- and site-scoped dashboards for the submitting clinician, MDT chair, specialist reviewers (hepatology, transplant and surgical, medical and radiation oncology, radiology, pathology), and the coordinator. Users see the steps and cases appropriate to their role and tenant (your FAQ Q44). Chair, specialist, and admin dashboard depth expands at Tier 2.

20. Patient journey tracking.

Yes. The core of LORiMDT is a longitudinal patient-journey data model that persists clinical decisions, clinical reasoning, and trigger state across the full arc of care, rather than point-in-time snapshots (your FAQ Q12). Source data stays in the EHR and LIS and is ingested via FHIR R4; LORiMDT maintains the persistent governance layer on top. A case is submitted, becomes read-only, is routed for review, produces a treatment plan, and re-opens through iterative clinical reassessment when conditions change.

21. Documentation and audit trail.

We build to the Tier 1 audit floor you set in your FAQ Q45: authentication events; PHI access (views and searches); case create, update, and delete with field-level change history; case state transitions; the clinical handshake events (ownership assignment and acknowledgment) with actor, timestamp, and captured intent; note generation and EHR write-back; escalation and soft-stop triggers and their resolution; configuration and permission changes; and integration transactions (FHIR pull and push) with success or failure. Entries are attributable via SSO identity, tamper-evident, and immutably timestamped. On the handshake (your FAQ Q23): for Tier 1 an attributable, timestamped, audit-logged acknowledgment is sufficient, and full 21 CFR Part 11 is not a Tier 1 requirement, but we architect the signature and audit layer to be Part 11-upgradeable without re-architecture for your trial-enabling roadmap.

2.4 Security and compliance

22. HIPAA compliance.

The platform runs in a US-based, HIPAA-eligible AWS environment with US data residency. We execute BAAs. PHI is encrypted, access is least-privilege and client-governed, and the audit trail (Q21) covers PHI access. Production PHI is processed only in the US environment.

23. Certifications.

We will be direct: ElectroCom does not currently hold SOC 2 Type II, HITRUST, or ISO 27001. What we do is build to the controls those frameworks specify (encryption, least-privilege access, audit logging, change management, and BAAs), and we will commit contractually to achieving SOC 2 Type II within 12 months of contract signing, with ISO 27001 for our delivery operations on a similar horizon. If a certification is a hard gate for you at pilot, tell us and we will scope the timeline and cost into the proposal directly rather than around it.

24. Data encryption.

Encryption at rest (AES-256, KMS-managed keys) and in transit (TLS 1.2 or higher). Field-level handling for the most sensitive PHI where warranted, with key rotation and access logging on the key store.

25. Access controls.

RBAC scoped by tenant and clinical role. Authentication via SSO and SMART on FHIR against each site's existing identity provider: SMART on FHIR for Epic sites, and SAML or OIDC federation to the site's enterprise IdP (Azure AD / Entra, Okta, or PingFederate) more generally. We maintain no separate password store (your FAQ Q41 and Q48). MFA is enforced at the IdP, with session management and re-authentication for sensitive actions. For patients, a stable internal patient identifier resolves site-specific MRNs via MPI-style matching, with human adjudication for uncertain matches, because in oncology a false merge is a patient-safety event (your FAQ Q31 and Q41).

26. Data residency.

Production PHI resides in US AWS regions. The application, its database, and its authentication endpoints deploy to your US-region AWS account, enforced by region restriction, network controls, and IAM policy (your FAQ Q48).

Disclosure, per your FAQ Q57 and Q59: ElectroCom's engineering team works from our Multan, Pakistan office, and our US entity and engagement accountability sit at our Celina, Texas headquarters. We state that plainly rather than implying a US-resident engineering bench we do not have.

The control that matters is not where engineers sit but who can reach production PHI:

- No standing human access to production. Deployment is via CI/CD from version control, and engineers hold no interactive production credentials.
- Development and testing run entirely against de-identified or synthetic data, never production PHI. That is the split your FAQ Q59 describes.
- Break-glass access is time-boxed, individually attributed, Lucas Health Tech-approved, fully session-logged, and reachable only through a jump path you control.
- Logs are PHI-scrubbed, so routine debugging never exposes patient data.

If you require that every individual capable of break-glass access be US-resident, we will staff that role US-side and price it. Tell us and we will scope it directly rather than around it.

2.5 Implementation approach and timeline

27. 5 to 6-month feasibility.

Yes, feasible for the Tier 1 scope on your timeline (kickoff August 19, 2026 to pilot go-live January 2027). We would rather deliver the Tier 1 accountability loop working end to end and trustworthy than a broad, shallow demo of every tier (your FAQ Q17). Proposed sequence:

- August to September: discovery with your clinical SMEs (Dr. Baker's group) and the pilot sites' Epic teams; ingest the validated field dictionary (212 case-submission fields and 74 reassessment fields, per your FAQ Q21); stand up the US AWS multi-tenant foundation, SSO, and the FHIR integration layer; begin Epic App Orchard registration.
- October: structured capture against the dictionary, the case-lifecycle state machine, the clinical handshake, and the Level 1 soft-stop; Epic FHIR R4 pull; begin the MDT-note write-back.
- October to November: integration testing at the pilot sites; note generation; hardening of pull and write-back.
- December: user acceptance testing and clinician training.
- January 2027: pilot go-live at 2 to 3 sites.

28. Methodology.

Agile, two-week sprints, with a fixed cadence of demo and retro and a clinical-review checkpoint each sprint against Dr. Baker's group. Because patient-safety logic is involved, rule and workflow changes go through the versioned, audit-logged governance path (Q15) rather than ad-hoc edits.

29. Configuration versus custom development.

For LORiMDT we estimate roughly 55 to 65 percent custom build (the clinician-facing application, the FHIR integration layer, the App Orchard note-write certification, and the longitudinal

governance model, which is a greenfield build per your FAQ Q15 and Q16) and 35 to 45 percent configuration (the workflow engine and rules engine configured to your defined workflows and guideline-anchored thresholds). We will refine this split during discovery.

30. Team composition.

A dedicated pod for the pilot. We name a core of six as fully dedicated and additional engineers as allocated, so the pilot is not exposed if another engagement lands.

Core, dedicated to the pilot:

- Arslan Noor, CTO. Engagement lead and solution architect (AI and full-stack, frontier-AI experience). United States.
- Haseeb Ilyas. Backend and FHIR integration. Multan, PK.
- Muhammad Husnain. Data engineering (longitudinal model, registry-quality capture). Multan, PK.
- Umar Farooq. Data science and AI (note generation, AI-assisted capture). Multan, PK.
- Bilal Arshad. DevOps and security (US AWS deployment, HIPAA controls, HA/DR). Multan, PK.
- Zain Goraya. Senior full-stack. Multan, PK.

Allocated as the work requires:

- Noman Aijaz. Full-stack. Multan, PK.
- Aadil Saeed, Shayan Ali. QA. Multan, PK.
- Clinical and product liaison, working directly with your BA/product owner and Dr. Baker's group.

Contracting entity and engagement accountability are US-based (Celina, TX). The engagement lead is US-based; all other named engineers work from our Pakistan office against de-identified data, per Q26.

31. Day 1 structured data capture.

Yes, and it is central to the design. The clinical workflow generates registry-quality structured data as a byproduct of the work itself, not through a separate data-entry step. Because the platform is greenfield (your FAQ Q16 and Q52) there is no legacy migration at Tier 1: from go-live, a clinician initiates a case and the platform pulls that patient's current EHR data via FHIR to populate it, with the three-tier capture model and required-field enforcement making each case decision-ready. Tier 1 supports the human Certified Tumor Registrar workflow with a structured, complete case summary that reduces manual abstraction; direct NAACCR XML and RCRS export is a later phase (your FAQ Q53).

2.6 Integration with Open Evidence

32. Open Evidence integration.

We understand Open Evidence embedded in the MDT workflow is a Tier 3 capability and not on the Tier 1 critical path (your FAQ Q10 and Q56), and we scope it accordingly. Our approach to embedding an NPI-credentialed clinical reference tool is to integrate rather than rebuild, launch it in context from the clinical workflow, and treat clinician credentialing and provisioning as a parallel operational workstream rather than a platform-build line item. We do not carry credential provisioning in the platform estimate (your FAQ Q56). We have integrated OAuth and credentialed third-party services in production before, and we would apply the same in-context launch and token-handling pattern here.

2.7 Pricing structure

33. Pricing model.

Because LORiMDT is intended to be Lucas Health Tech-owned IP (your FAQ Q18 and Q62), we price predominantly as one-time build and implementation, with a modest recurring component limited to hosting, support, and maintenance, and no application-license fee that would amount to renting the product back to you (your FAQ Q63). The three components are separated below.

34. Implementation costs (Tier 1, 2 to 3-site MVP).

One-time build, implementation, and professional services: **\$385,000**. This covers the multi-tenant US AWS foundation; the FHIR integration layer and the App Orchard note-write certification; structured capture to the field dictionary; the case lifecycle, clinical handshake, and Level 1 soft-stop; automated note generation; essential reporting (case volume, completion rate, time-to-decision); RBAC, SSO, and audit logging; and HIPAA controls.

35. Ongoing costs (annual, post-go-live).

- Hosting on your AWS (HA/DR, multi-region): pass-through at cloud cost plus platform management of roughly \$30,000 per year.
- Support and maintenance: roughly \$66,000 per year, about 17 percent of the build, covering monitoring, updates, bug fixes, and a defined support SLA.
- Recurring software or application license: \$0. The application IP is yours, so there is no application-rental license. This answers the concern your FAQ Q63 raises directly.

36. Startup-friendly terms.

Yes. We propose milestone-based payments aligned to your funding-close dates, with a committed minimum funding the early sprints so no vendor is asked to carry the full financial risk (your FAQ Q19). Proposed schedule against the \$385,000 one-time build:

Milestone	Timing	Share	Amount
Kickoff (committed minimum)	August 2026	20%	\$77,000
Discovery and foundation complete	End September 2026	25%	\$96,250
Integration live	End October 2026	25%	\$96,250
UAT complete	December 2026	20%	\$77,000
Pilot go-live	January 2027	10%	\$38,500

The first three milestones (70%, roughly \$269,500) fund the August to October build and are the committed floor. We can align the final two milestones (30%, roughly \$115,500) to your capital close, so we carry that timing with you rather than requiring it on a fixed date. We would rather structure the shared risk as a modest payment deferral than as equity or a pure success fee, since a deferral keeps the engagement simple and our team fully staffed. If a small go-live success bonus fits your model better, we are open to it, but the deferral is our recommendation.

Before you commit to the full build, we also offer a fixed-fee discovery sprint of two to three weeks at \$20,000: a FHIR connectivity spike against a sandbox, the field-dictionary data model stood up, and one working clinical-handshake path. It is not additive to the schedule. It is the first two to three weeks of the Q27 sequence (your August discovery window), so it does not move the January go-live. If you proceed to the full build, the \$20,000 is credited in full against the kickoff milestone, so it adds nothing to the \$385,000. If you walk away at the end, you keep the artifacts and owe nothing further. Rather than ask you to take our references on faith, this lets you watch us work against your actual problem for three weeks before the larger commitment.

2.8 Budget and phasing approach

37. Our recommendation.

Tier 1 (Bootstrap MVP), with a pre-agreed phased upgrade to Tier 2 on your August funding close. The reasoning:

- Your FAQ Q17 is explicit that a successful pilot is defined by the Tier 1 accountability loop working end to end (structured capture to the dictionary, the clinical handshake, bidirectional Epic FHIR pull and MDT-note write-back, automated note generation, and the

Level 1 soft-stop), not by demonstrating Level 2 and 3 monitoring and escalation. Tier 1 is the honest scope for a first pilot at 2 to 3 sites.

- Your funding closes in August, and you invited a phased approach (RFI 2.8, your FAQ Q61). Committing you to Tier 2 or 3 scope before the close would invert the shared-risk structure you asked for.
- We would rather earn Tier 2 by delivering Tier 1 well than win a larger number and underdeliver on breadth.

We are not bidding low to win. We believe Tier 1 is the right scope, and Q39 states exactly what we would add at Tier 2 so you can judge that for yourself.

38. Detailed proposal (Tier 1).

- Scope and deliverables: as listed in Q34, mapped one to one against your Tier 1 required-features list (MDT case presentation with structured capture; the Digital Handshake; basic Epic FHIR R4 pull and MDT-note push; soft-stop validation; automated note generation; essential reporting; RBAC, SSO, and audit; HIPAA compliance).
- Timeline and milestones: per Q27 (kickoff August 19 to go-live January 2027), with milestone gates at discovery-complete (September), integration-live (October), and UAT-complete (December).
- Team and roles: per Q30.
- Cost breakdown: per Q34 and Q35.
- Payment terms: milestone-based and funding-aligned (Q36).
- Assumptions and exclusions: greenfield, so no legacy data migration at Tier 1 (your FAQ Q52); pilot sites and their Epic instance details confirmed during discovery, with Lucas Health Tech brokering site technical-team access (your FAQ Q50); Open Evidence, NLP-from-narrative, ePRO, multi-cancer expansion, and direct registry export are out of Tier 1 scope; clinician credential provisioning is an operational workstream, not a platform line item.

39. Alternative tier proposal (Tier 2, brief).

If your close lands early and you want the larger cohort from the start, Tier 2 (roughly \$600,000, 3 to 5 sites) adds the full three-level escalation architecture (Level 2 execution monitoring on the 14, 21, and 28-day configurable default cadence, noting these are defaults and not guideline-mandated intervals, per your FAQ Q54; and Level 3 reassessment enforcement), ongoing bidirectional FHIR monitoring, advanced dashboards, broader multi-site support, and HITRUST-aligned security hardening. We would scope this as an increment on the Tier 1 build, not a rebuild.

40. Phasing plan.

We build Tier 1 on a Tier 2-ready architecture. Multi-tenancy, the configurable rules and escalation engine, and the audit layer are all present from day one (your FAQ Q39 and Q23), so Tier 2 features switch on as configuration and incremental build rather than re-architecture. Concretely: start Tier 1 in August against the committed minimum; when your capital close completes, execute a pre-agreed Tier 2 change order that adds Level 2 and 3 escalation, the advanced dashboards, and the additional sites, with payments re-aligned to the new milestones.

2.9 References

41. References.

Our most comparable reference is a live, multi-hospital clinical platform our team built and runs: structured clinical workflow on real patient data, integration with EHR and legacy systems, and continuous operation across hospitals in Australia and the UK. That client is under confidentiality, so we do not name them in writing. We introduce them directly as a reference at the demo stage, where your process schedules reference checks (August 5 to 7). This is our closest match to LORiMDT in clinical scope, and we would rather have you speak with them under NDA than describe them secondhand here.

Two named healthcare clients you can contact now:

1. **Sono Stat Imaging (Dallas-Fort Worth, Texas, US).** A US healthcare provider delivering physician-referred diagnostic ultrasound to patients: cardiac echocardiography, carotid and vascular Doppler, and abdominal and thyroid imaging, performed by credentialed sonographers and billed through insurance. ElectroCom built and runs the digital platform and operational systems behind their imaging service, the patient-facing and clinical-operations layer they rely on day to day. Contact: Iftikhar Saeed Khan, +1 210-422-0737, sonostatimaging@gmail.com.

2. **The Islamia University of Bahawalpur (IUB), Faculty of Medicine and Allied Health Sciences.** A university medical faculty with a teaching-hospital partner (Nawab Sir Sadiq Muhammad Khan Abbasi Hospital). ElectroCom built their clinical patient-management system. Contact: IT Director, +92 300 9635022.

For completeness, we also build data-integrity and manufacturing-accuracy systems for regulated health-product manufacturers (CoreSpeed and PharmaBuilt; contact John Linss, +1 770-630-9956, jlinss@corespeed.com). We list these as systems-engineering work, not clinical-workflow references, so each is weighed for what it is.

Appendix A. Case study: legacy clinical platform modernization (client anonymized)

Engagement. Modernization of a legacy desktop clinical application into a production cloud web platform, delivered by ElectroCom's engineering team.

Starting point. The client ran a desktop application that had served hospital clinical and diagnostic-lab workflows for over a decade. It was stable but aging: the deployment model blocked multi-location and remote access, the interface had fallen behind current clinical-software expectations, and the architecture could not support the client's expansion plans.

What we built. Over roughly 20 months we rebuilt the platform module by module into a modern React and TypeScript cloud application, while the existing system stayed live for hospital users. We preserved the backend and clinical-data integrity while replacing the frontend architecture, introduced a design system for long-term maintainability, and validated each module in real operational use before moving to the next. Modules delivered included quality control, patient records, clinical reporting and diagnostic-test workflows, appointment scheduling and templates, and staff rostering.

Result. A cloud-based, vendor-neutral clinical platform now live across hospitals in Australia and the UK, on an architecture built for cross-market expansion. Delivery was incremental and never took the clinical service offline.

Relevance to LORiMDT. The same terrain: real clinical data, multi-hospital operation, modernization under live operational constraints, and a data model built to be vendor-neutral and standards-ready. The stack (React, TypeScript, containerized services, cloud deployment) is the one we propose for LORiMDT.

Full detail, including the client, is available at the demo stage.